

[[2022-04-06]]

We would like to begin by thanking the reviewers for their thoughtful comments. One theme across the comments is the need for more thorough experimental evaluation of our system. While we believe more evaluation would be useful, we chose our benchmarks based on their frequency within applications that use RAJA. We would also like to clarify how our approach works with regards to loop bounds. While we know the loop bounds at runtime and use those values in our modeling, these values do not need to be compile-time constant.

We conclude with a few requests for the reviewers. First, please provide specific related works that you recommend we consider and cite. Second, please indicate which limitations you consider most significant/pressing (i.e., please indicate where you feel we should focus effort in extending the approach). Finally, please suggest applications (or computational domains) that we should consider for further evaluation. Thank you again for your time and reviews.